

# Data Mining - Lab - 2

## Numpy & Perform Data Exploration with Pandas

Mohil Parmar

23010101192

---

### Numpy

1. NumPy (Numerical Python) is a powerful open-source library in Python used for numerical and scientific computing.
2. It provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on them efficiently.
3. NumPy is highly optimized and written in C, making it much faster than using regular Python lists for numerical operations.
4. It serves as the foundation for many other Python libraries in data science and machine learning, like pandas, TensorFlow, and scikit-learn.
5. With features like broadcasting, vectorization, and integration with C/C++ code, NumPy allows for cleaner and faster code in numerical computations.

### Step 1. Import the Numpy library

```
In [5]: import numpy as np
```

### Step 2. Create a 1D array of numbers

```
In [11]: arr = np.arange(10)
arr
```

```
Out[11]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [13]: arr = np.arange(10,21)
arr
```

```
Out[13]: array([10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20])
```

```
In [15]: arr1 = np.array([22,33,44,55])
arr1
```

```
Out[15]: array([22, 33, 44, 55])
```

### Step 3. Reshape 1D to 2D Array

```
In [19]: arr2 = np.arange(20).reshape(4,5)
arr2
```

```
Out[19]: array([[ 0,  1,  2,  3,  4],
               [ 5,  6,  7,  8,  9],
               [10, 11, 12, 13, 14],
               [15, 16, 17, 18, 19]])
```

### Step 4. Create a Linspace array

```
In [21]: np.linspace(14,18)
# by default 50 decimal
```

```
Out[21]: array([14.          , 14.08163265, 14.16326531, 14.24489796, 14.32653061,
               14.40816327, 14.48979592, 14.57142857, 14.65306122, 14.73469388,
               14.81632653, 14.89795918, 14.97959184, 15.06122449, 15.14285714,
               15.2244898 , 15.30612245, 15.3877551 , 15.46938776, 15.55102041,
               15.63265306, 15.71428571, 15.79591837, 15.87755102, 15.95918367,
               16.04081633, 16.12244898, 16.20408163, 16.28571429, 16.36734694,
               16.44897959, 16.53061224, 16.6122449 , 16.69387755, 16.7755102 ,
               16.85714286, 16.93877551, 17.02040816, 17.10204082, 17.18367347,
               17.26530612, 17.34693878, 17.42857143, 17.51020408, 17.59183673,
               17.67346939, 17.75510204, 17.83673469, 17.91836735, 18.          ])
```

```
In [23]: np.linspace(14,18,20)
        # in third argument use specify length of number
```

```
Out[23]: array([14.          , 14.21052632, 14.42105263, 14.63157895, 14.84210526,
               15.05263158, 15.26315789, 15.47368421, 15.68421053, 15.89473684,
               16.10526316, 16.31578947, 16.52631579, 16.73684211, 16.94736842,
               17.15789474, 17.36842105, 17.57894737, 17.78947368, 18.          ])
```

## Step 5. Create a Random Numbered Array

```
In [25]: arr4 = np.random.rand(6)
        arr4
```

```
Out[25]: array([0.76539575, 0.54073109, 0.61656565, 0.50856795, 0.9464426 ,
               0.5109816 ])
```

```
In [27]: arr5=np.random.rand(6,5)
        arr5
        # for 2D array use give second argument
```

```
Out[27]: array([[0.46938984, 0.75824529, 0.79401908, 0.667392 , 0.86403358],
               [0.06399315, 0.35781957, 0.39032216, 0.40778384, 0.93098397],
               [0.59229555, 0.3266491 , 0.20554348, 0.171566 , 0.30838983],
               [0.8795907 , 0.31034884, 0.55286617, 0.58000734, 0.36457935],
               [0.50805045, 0.40600399, 0.88392335, 0.17804145, 0.09062223],
               [0.83127196, 0.62684204, 0.8553937 , 0.47073148, 0.75429648]])
```

## Step 6. Create a Random Integer Array

```
In [29]: np.random.randint(30,60)
```

```
Out[29]: 55
```

```
In [31]: np.random.randint(20,40,size=5)
```

```
Out[31]: array([34, 25, 22, 36, 34])
```

```
In [33]: np.random.randint(30,50,size=(3,4))
```

```
Out[33]: array([[38, 37, 38, 44],  
               [31, 47, 44, 45],  
               [36, 32, 45, 47]])
```

## Step 7. Create a 1D Array and get Max,Min,ArgMax,ArgMin

```
In [53]: arr6 = np.random.randint(10,40,size=10)  
print(arr6)  
# max(arr6)  
arr6.max()
```

```
[21 27 24 14 37 25 21 10 22 22]
```

```
Out[53]: 37
```

```
In [55]: arr6.min()
```

```
Out[55]: 10
```

```
In [57]: arr6.argmax()
```

```
Out[57]: 4
```

```
In [59]: arr6.argmin()
```

```
Out[59]: 7
```

## Step 8. Indexing in 1D Array

```
In [63]: arr7 = np.random.randint(10,50,size=7)
arr7
```

```
Out[63]: array([28, 33, 47, 35, 48, 13, 41])
```

```
In [67]: arr7[5]
```

```
Out[67]: 13
```

```
In [69]: arr7[2:6]
```

```
Out[69]: array([47, 35, 48, 13])
```

```
In [77]: arr7.dtype
print(type(arr7))
```

```
<class 'numpy.ndarray'>
```

## Step 9. Indexing in 2D Array

```
In [89]: arr8 = np.random.randint(10,40,size=(5,5))
arr8
```

```
Out[89]: array([[23, 17, 23, 10, 37],
                [15, 14, 14, 14, 39],
                [33, 34, 29, 34, 15],
                [11, 38, 32, 26, 24],
                [17, 31, 39, 18, 14]])
```

```
In [91]: arr8[3]
```

```
Out[91]: array([11, 38, 32, 26, 24])
```

```
In [95]: arr8[0:3]
```

```
Out[95]: array([[23, 17, 23, 10, 37],
               [15, 14, 14, 14, 39],
               [33, 34, 29, 34, 15]])
```

```
In [109... arr8[3][3]
```

```
Out[109... 26
```

## Step 10. Conditional Selection

```
In [121... arr10 = np.random.randint(20,40,size=10)
a1=[i for i in arr10 if i>25]
a1
```

```
Out[121... [35, 29, 26, 37, 27, 26]
```

```
In [123... arr10[arr10>25]
```

```
Out[123... array([35, 29, 26, 37, 27, 26])
```

```
In [127... arr10[(arr10>25) & (arr10<30)]
```

```
Out[127... array([29, 26, 27, 26])
```

🔥 You did it! 10 exercises down — you're on fire! 🔥

## Pandas

### Step 1. Import the necessary libraries

```
In [131... import pandas as pd
```

### Step 2. Import the dataset from this [address](#).

### Step 3. Assign it to a variable called users and use the 'user\_id' as index

```
In [148... users = pd.read_csv("https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user" , sep="|" , index_col="user_id")
users
```

```
Out[148...      age  gender  occupation  zip_code
user_id
1      24     M    technician  85711
2      53     F         other  94043
3      23     M         writer  32067
4      24     M    technician  43537
5      33     F         other  15213
...     ...     ...         ...      ...
939    26     F         student  33319
940    32     M  administrator  02215
941    20     M         student  97229
942    48     F         librarian  78209
943    22     M         student  77841
```

943 rows × 4 columns

### Step 4. See the first 25 entries

```
In [150... users.head(25)
```

Out[150...

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213
6	42	M	executive	98101
7	57	M	administrator	91344
8	36	M	administrator	05201
9	29	M	student	01002
10	53	M	lawyer	90703
11	39	F	other	30329
12	28	F	other	06405
13	47	M	educator	29206
14	45	M	scientist	55106
15	49	F	educator	97301
16	21	M	entertainment	10309
17	30	M	programmer	06355
18	35	F	other	37212
19	40	M	librarian	02138
20	42	F	homemaker	95660
21	26	M	writer	30068



	age	gender	occupation	zip_code
user_id				
22	25	M	writer	40206
23	30	F	artist	48197
24	21	F	artist	94533
25	39	M	engineer	55107

## Step 5. See the last 10 entries

In [152... `users.tail(10)`

Out[152...

	age	gender	occupation	zip_code
user_id				
934	61	M	engineer	22902
935	42	M	doctor	66221
936	24	M	other	32789
937	48	M	educator	98072
938	38	F	technician	55038
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

## Step 6. What is the number of observations in the dataset?

```
In [162... users.shape[0]
```

```
Out[162... 943
```

## Step 7. What is the number of columns in the dataset?

```
In [170... users.shape[1]
```

```
Out[170... 4
```

## Step 8. Print the name of all the columns.

```
In [174... users.columns
```

```
Out[174... Index(['age', 'gender', 'occupation', 'zip_code'], dtype='object')
```

## Step 9. How is the dataset indexed?

```
In [176... users.index
```

```
Out[176... Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,  
      ...  
      934, 935, 936, 937, 938, 939, 940, 941, 942, 943],  
      dtype='int64', name='user_id', length=943)
```

## Step 10. What is the data type of each column?

```
In [178... users.dtypes
```

```
Out[178... age          int64
gender       object
occupation   object
zip_code     object
dtype: object
```

## Step 11. Print only the occupation column

```
In [184... users["occupation"]
```

```
Out[184... user_id
1          technician
2              other
3              writer
4          technician
5              other
...
939         student
940 administrator
941         student
942         librarian
943         student
Name: occupation, Length: 943, dtype: object
```

## Step 12. How many different occupations are in this dataset?

```
In [190... users["occupation"].nunique()
# users.occupation.nunique()
```

```
Out[190... 21
```

```
In [192... users.occupation.unique()
```

```
Out[192... array(['technician', 'other', 'writer', 'executive', 'administrator',
      'student', 'lawyer', 'educator', 'scientist', 'entertainment',
      'programmer', 'librarian', 'homemaker', 'artist', 'engineer',
      'marketing', 'none', 'healthcare', 'retired', 'salesman', 'doctor'],
      dtype=object)
```

### Step 13. What is the most frequent occupation?

```
In [196...] users["occupation"].value_counts()
```

```
Out[196...] occupation
student      196
other        105
educator     95
administrator 79
engineer     67
programmer   66
librarian    51
writer       45
executive    32
scientist    31
artist       28
technician   27
marketing    26
entertainment 18
healthcare   16
retired      14
lawyer       12
salesman     12
none         9
homemaker    7
doctor       7
Name: count, dtype: int64
```

```
In [198...] users["occupation"].value_counts().head(1)
```

```
Out[198...] occupation
student      196
Name: count, dtype: int64
```

```
In [200...] users["occupation"].value_counts().idxmax()
```

```
Out[200...] 'student'
```

## Step 14. Summarize the DataFrame.

```
In [202... users.describe()
```

```
Out[202...  
      age  
count  943.000000  
mean   34.051962  
std    12.192740  
min     7.000000  
25%    25.000000  
50%    31.000000  
75%    43.000000  
max    73.000000
```

## Step 15. Summarize all the columns

```
In [214... users.describe(include="all")
```

Out[214...

	age	gender	occupation	zip_code
<b>count</b>	943.000000	943	943	943
<b>unique</b>	NaN	2	21	795
<b>top</b>	NaN	M	student	55414
<b>freq</b>	NaN	670	196	9
<b>mean</b>	34.051962	NaN	NaN	NaN
<b>std</b>	12.192740	NaN	NaN	NaN
<b>min</b>	7.000000	NaN	NaN	NaN
<b>25%</b>	25.000000	NaN	NaN	NaN
<b>50%</b>	31.000000	NaN	NaN	NaN
<b>75%</b>	43.000000	NaN	NaN	NaN
<b>max</b>	73.000000	NaN	NaN	NaN

## Step 16. Summarize only the occupation column

In [222...

```
users["occupation"].describe()
```

Out[222...

```
count      943
unique      21
top        student
freq        196
Name: occupation, dtype: object
```

## Step 17. What is the mean age of users?

In [224...

```
users["age"].mean()
```

Out[224...

```
34.05196182396607
```

## Step 18. What is the age with least occurrence?

```
In [226... users.age.value_counts().tail()
```

```
Out[226... age
7      1
66     1
11     1
10     1
73     1
Name: count, dtype: int64
```

You're not just learning, you're mastering it. Keep aiming higher! 🚀